

# Data Science Project 3

- Student Name: Natalya Doris
- Student Pace: Flex / 40 weeks
- Scheduled Project Review Date / Time: Thurs, Sept 16 / 12pm
- Instructor Name: Abhineet Kulkarni

## Setup, EDA, Preprocessing

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from datetime import datetime

from sklearn.model_selection import train_test_split, GridSearchCV, \
cross_val_score
from sklearn.impute import SimpleImputer
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, \
BaggingClassifier, ExtraTreesClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, f1_score, recall_score, \
precision_score, confusion_matrix, classification_report, roc_curve, auc, \
average_precision_score
from sklearn.preprocessing import StandardScaler, LabelEncoder, \
MinMaxScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE

import warnings
warnings.filterwarnings('ignore')
```

## Load in Data

```
In [2]: df = pd.read_csv('data/investments_VC.csv', encoding = "ISO-8859-1")
```

## Basic Cleaning

```
In [3]: # getting rid of extra spaces in market and funding_total_usd
df = df.rename(columns={' market ': 'market',
                        ' funding_total_usd ': 'funding_total_usd'})
```

Dropping irrelevant columns:

```
In [4]: df = df.drop(columns=['permalink', 'homepage_url', 'category_list',  
                             'founded_quarter', 'post_ipo_equity',  
                             'post_ipo_debt', 'secondary_market'],  
                  axis=1)
```

```
In [5]: # converting to float  
df['funding_total_usd'][~df['funding_total_usd'].isnull()] = \  
    [float(num.replace(" ", "0").replace(",", "0").replace("-", "0")) \  
     for num in df['funding_total_usd'][~df['funding_total_usd'].isnull()]]
```

Dropping duplicates, if any:

```
In [6]: df = df.drop_duplicates()
```

```
In [7]: df['status'].value_counts()
```

```
Out[7]: operating      41829  
        acquired       3692  
        closed         2603  
        Name: status, dtype: int64
```

## Exploratory Analysis - Full Dataset

```
In [8]: data = df
```

```
In [9]: print(data.info())  
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 49439 entries, 0 to 49438
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   name                                  49437 non-null  object
1   market                               45470 non-null  object
2   funding_total_usd                    49438 non-null  object
3   status                               48124 non-null  object
4   country_code                         44165 non-null  object
5   state_code                           30161 non-null  object
6   region                               44165 non-null  object
7   city                                 43322 non-null  object
8   funding_rounds                       49438 non-null  float64
9   founded_at                           38554 non-null  object
10  founded_month                         38482 non-null  object
11  founded_year                         38482 non-null  float64
12  first_funding_at                     49438 non-null  object
13  last_funding_at                      49438 non-null  object
14  seed                                 49438 non-null  float64
15  venture                              49438 non-null  float64
16  equity_crowdfunding                  49438 non-null  float64
17  undisclosed                          49438 non-null  float64
18  convertible_note                     49438 non-null  float64
19  debt_financing                       49438 non-null  float64
20  angel                                49438 non-null  float64
21  grant                                49438 non-null  float64
22  private_equity                       49438 non-null  float64
23  product_crowdfunding                 49438 non-null  float64
24  round_A                             49438 non-null  float64
25  round_B                             49438 non-null  float64
26  round_C                             49438 non-null  float64
27  round_D                             49438 non-null  float64
28  round_E                             49438 non-null  float64
29  round_F                             49438 non-null  float64
30  round_G                             49438 non-null  float64
31  round_H                             49438 non-null  float64
dtypes: float64(20), object(12)
memory usage: 12.4+ MB
None
```

Out[9]:

	name	market	funding_total_usd	status	country_code	state_code	region
--	------	--------	-------------------	--------	--------------	------------	--------

0	#waywire	News	1.705e+09	acquired	USA	NY	New York City
1	&TV Communications	Games	4e+09	operating	USA	CA	Los Angeles
2	'Rock' Your Paper	Publishing	4e+06	operating	EST	NaN	Tallinn
3	(In)Touch Network	Electronics	1.5e+09	operating	GBR	NaN	London
4	-R- Ranch and Mine	Tourism	6e+06	operating	USA	TX	Dallas

5 rows x 32 columns

In [10]: `df.describe()`

Out[10]:

	funding_rounds	founded_year	seed	venture	equity_crowdfunding	un
<b>count</b>	49438.000000	38482.000000	4.943800e+04	4.943800e+04	4.943800e+04	4.94
<b>mean</b>	1.696205	2007.359129	2.173215e+05	7.501051e+06	6.163322e+03	1.30
<b>std</b>	1.294213	7.579203	1.056985e+06	2.847112e+07	1.999048e+05	2.98
<b>min</b>	1.000000	1902.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.00
<b>25%</b>	1.000000	2006.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.00
<b>50%</b>	1.000000	2010.000000	0.000000e+00	0.000000e+00	0.000000e+00	0.00
<b>75%</b>	2.000000	2012.000000	2.500000e+04	5.000000e+06	0.000000e+00	0.00
<b>max</b>	18.000000	2014.000000	1.300000e+08	2.351000e+09	2.500000e+07	2.92

In [11]: `data.isnull().sum()`

Out[11]:

name	2
market	3969
funding_total_usd	1
status	1315
country_code	5274
state_code	19278
region	5274
city	6117
funding_rounds	1
founded_at	10885
founded_month	10957
founded_year	10957
first_funding_at	1
last_funding_at	1
seed	1
venture	1
equity_crowdfunding	1
undisclosed	1
convertible_note	1
debt_financing	1
angel	1
grant	1
private_equity	1
product_crowdfunding	1
round_A	1
round_B	1
round_C	1
round_D	1
round_E	1
round_F	1
round_G	1
round_H	1
dtype:	int64

In [12]: `data['founded_at'].head()`

```
Out[12]: 0    2012-06-01
          1         NaN
          2    2012-10-26
          3    2011-04-01
          4    2014-01-01
          Name: founded_at, dtype: object
```

```
In [13]: data['region'].value_counts()
```

```
Out[13]: SF Bay Area      6804
          New York City   2577
          Boston         1837
          London         1588
          Los Angeles    1389
          ...
          Stirling        1
          Igualada       1
          MUS - Other     1
          Cheadle Hulme  1
          Reigate        1
          Name: region, Length: 1089, dtype: int64
```

```
In [14]: data['city'].value_counts()
```

```
Out[14]: San Francisco    2615
          New York         2334
          London          1257
          Palo Alto       597
          Austin          583
          ...
          Tel Mond        1
          Ra'ananna       1
          Southwark       1
          Benson          1
          Alvorada        1
          Name: city, Length: 4188, dtype: int64
```

```
In [15]: data['state_code'].value_counts()
```

```
Out[15]: CA      9917
          NY      2914
          MA      1969
          TX      1466
          WA       974
          ...
          MB       13
          AK       12
          NB        8
          SK        4
          PE        2
          Name: state_code, Length: 61, dtype: int64
```

```
In [16]: data['country_code'].value_counts()
```


```
Out[16]: USA      28793
        GBR      2642
        CAN      1405
        CHN      1239
        DEU       968
        ...
        SOM        1
        OMN        1
        CIV        1
        TTO        1
        MOZ        1
        Name: country_code, Length: 115, dtype: int64
```

```
In [17]: # inspecting market feature - lots of catgories
print("Number of unique markets: ", len(set(data['market'])))
print("Markets with more than 200 companies: ",
      str(sum(data['market'].value_counts() >= 200)))
```

```
Number of unique markets: 754
Markets with more than 200 companies: 39
```

```
In [18]: fig, ax = plt.subplots(figsize = (10,10))
data['market'].value_counts()[ :30].plot(kind='barh')
```

```
Out[18]: <AxesSubplot:>
```



```
In [19]: # same deal with region
print("Number of unique regions: ", len(set(data['region'])))
```


```
print("Regions with more than 200 companies: ",
      str(sum(data['region'].value_counts() >= 200)))
```

Number of unique regions: 1090

Regions with more than 200 companies: 44


```
In [20]: fig, ax = plt.subplots(figsize = (10,10))
data['region'].value_counts()[:30].plot(kind='barh')
```

Out[20]: <AxesSubplot:>



```
In [21]: data.hist('equity_crowdfunding', bins=40)
# there are some very skewed columns in the funding area
```

Out[21]: array([[<AxesSubplot:title={'center': 'equity\_crowdfunding'}>]],  
dtype=object)



## Preprocessing & Feature Engineering

In [22]: `data.columns`

Out[22]: Index(['name', 'market', 'funding\_total\_usd', 'status', 'country\_code', 'state\_code', 'region', 'city', 'funding\_rounds', 'founded\_at', 'founded\_month', 'founded\_year', 'first\_funding\_at', 'last\_funding\_at', 'seed', 'venture', 'equity\_crowdfunding', 'undisclosed', 'convertible\_note', 'debt\_financing', 'angel', 'grant', 'private\_equity', 'product\_crowdfunding', 'round\_A', 'round\_B', 'round\_C', 'round\_D', 'round\_E', 'round\_F', 'round\_G', 'round\_H'], dtype='object')

## Missing Values

- I am removing rows with missing status or company name.
- I am also dropping rows with a missing founded\_year because this feature has high importance in many of the models & thus would prefer not to impute at risk of skewing the model. I tried both dropping and keeping the missing founded\_years (imputed). The proportion of companies in each status category & model results are fairly similar with & without imputing the median of year so I am comfortable dropping these rows.

In [23]: `data = data.dropna(subset=['status', 'name', 'founded_year'])`  
`data.isnull().sum()`



```
Out[23]: name                                0
market                                1801
funding_total_usd                     0
status                                0
country_code                          2936
state_code                           13332
region                               2936
city                                 3357
funding_rounds                        0
founded_at                           0
founded_month                         0
founded_year                         0
first_funding_at                     0
last_funding_at                      0
seed                                  0
venture                              0
equity_crowdfunding                   0
undisclosed                          0
convertible_note                     0
debt_financing                       0
angel                                0
grant                                 0
private_equity                       0
product_crowdfunding                  0
round_A                              0
round_B                              0
round_C                              0
round_D                              0
round_E                              0
round_F                              0
round_G                              0
round_H                              0
dtype: int64
```

```
In [24]: len(data)
```

```
Out[24]: 37563
```

```
In [25]: # filling categoricals
data = data.fillna(value={'market': 'other', 'country_code': 'other',
                          'region': 'other', 'city': 'other',
                          'state_code': 'other'})
```

```
In [26]: # distribution of non NA years
plt.hist(data['founded_year'][~data['founded_year'].isnull()], bins=30)
plt.axvline(x=np.nanmean(data['founded_year']), color='r')
plt.axvline(x=np.nanmedian(data['founded_year']), color='y')


p10 = np.percentile(data['founded_year'][~data['founded_year'].isnull()], 10)
print("10th percentile of founded_year: ", p10)

#10th percentile year is 2000, we will subset for startups founded on or after
print("Num Rows before 2000: ", len(data[data['founded_year'] < 2000]))
data = data[data['founded_year'] >= 2000].reset_index(drop=True)
len(data)
```

```
10th percentile of founded_year: 2000.0
```

```
Num Rows before 2000: 3518
```

```
Out[26]: 34045
```



```
In [27]: len(df[(df['founded_year'] < 2000) & (df['status'] != 'operating')])
```

```
Out[27]: 838
```

Date data - convert to datetime:

```
In [28]: data['founded_at'] = [datetime.strptime(day, '%Y-%m-%d') for day \
                             in data['founded_at'][~data['founded_at'].isnull()]]
data['first_funding_at'] = [datetime.strptime(day, '%Y-%m-%d').date() for day \
                             in data['first_funding_at']]
data['last_funding_at'] = [datetime.strptime(day, '%Y-%m-%d').date() for day \
                             in data['last_funding_at']]
data['founded_month'] = [datetime.strptime(mth, "%Y-%m").month for mth in data[
```

Fill year & founded\_at with simple imputer (only if not removing NaN rows). Since the data is skewed (see plot above), we will use median rather than mean.

```
In [29]: # uncomment the below to impute founded_year with mean, only if we are not
         # dropping nans for this feature

         #imp_median = SimpleImputer(missing_values=np.nan, strategy='median')

         #data['founded_year'] = imp_median.fit_transform(data[['founded_year']])
         #data['founded_at'] = imp_median.fit_transform(data[['founded_at']])
```

```
In [30]: data['founded_at'] = pd.to_datetime(data['founded_at'])
         data['founded_year'] = [day.year for day in data['founded_at']]
```

```
In [31]: data['founded_at'].value_counts()
```

```
Out[31]: 2012-01-01    2100
         2011-01-01    2096
         2010-01-01    1810
         2009-01-01    1561
         2013-01-01    1535
         ...
         2004-04-25      1
         2009-08-12      1
         2014-08-04      1
         2002-08-02      1
         2010-11-06      1
         Name: founded_at, Length: 2935, dtype: int64
```

```
In [32]: data.isna().sum()
```

```
Out[32]: name                0
         market              0
         funding_total_usd   0
         status              0
         country_code        0
         state_code          0
         region              0
         city                0
         funding_rounds      0
         founded_at          0
         founded_month       0
         founded_year        0
         first_funding_at    0
         last_funding_at     0
         seed                0
         venture             0
         equity_crowdfunding  0
         undisclosed         0
         convertible_note    0
         debt_financing      0
         angel               0
         grant               0
         private_equity      0
         product_crowdfunding 0
         round_A             0
         round_B             0
         round_C             0
         round_D             0
         round_E             0
         round_F             0
         round_G             0
         round_H             0
         dtype: int64
```

Basic data cleaning:

```
In [33]: # getting rid of extra spaces in market, city, state code, region
         data['market'] = [x.strip() for x in data['market']]
         data['country_code'] = [x.strip() for x in data['country_code']]
         data['state_code'] = [x.strip() for x in data['state_code']]
         data['region'] = [x.strip() for x in data['region']]
         data['city'] = [x.strip() for x in data['city']]
```

```
In [34]: data.dtypes
```

```
Out[34]: name                object
market                object
funding_total_usd     object
status                object
country_code          object
state_code            object
region                object
city                  object
funding_rounds        float64
founded_at            datetime64[ns]
founded_month         int64
founded_year          int64
first_funding_at      object
last_funding_at       object
seed                  float64
venture                float64
equity_crowdfunding    float64
undisclosed            float64
convertible_note       float64
debt_financing         float64
angel                  float64
grant                  float64
private_equity         float64
product_crowdfunding   float64
round_A                float64
round_B                float64
round_C                float64
round_D                float64
round_E                float64
round_F                float64
round_G                float64
round_H                float64
dtype: object
```

```
In [35]: # need to convert this data type to integer
data['funding_total_usd'].value_counts()
```

```
Out[35]: 0.000000e+00    5632
1.000000e+09      627
1.000000e+08      582
5.000000e+08      573
4.000000e+06      466
...
3.012099e+11       1
1.045077e+08       1
1.000108e+12       1
1.609205e+09       1
2.068042e+11       1
Name: funding_total_usd, Length: 10560, dtype: int64
```

```
In [36]: # data['funding_total_usd'] = [float(num.replace(" ", "0").replace(",","0").\
#                                     replace("-", "0"))
#                                     for num in data['funding_total_usd']]
# data['funding_total_usd'].dtypes
```

## Feature Engineering

```
In [37]: # creating column that labels country as domestic or international
data['international'] = [0 if country=='USA' else 1 for country in data['country']]
```

```
In [38]: # creating temporary columns to aid in calculation of time to first funding

data['founded_at_temp'] = [day.date() for day in data['founded_at']]
data['founded_at_temp'] = pd.to_datetime(data['founded_at_temp'],
                                         format = '%Y-%m-%d')

data['first_funding_at_temp'] = pd.to_datetime(data['first_funding_at'],
                                              format = '%Y-%m-%d',
                                              errors='coerce')
data['last_funding_at_temp'] = pd.to_datetime(data['last_funding_at'],
                                              format = '%Y-%m-%d',
                                              errors='coerce')

data['time_to_first_funding'] = (data['first_funding_at_temp'] - \
                                data['founded_at_temp']) / pd.Timedelta(days=365)

data['time_first_to_last_funding'] = (data['last_funding_at_temp'] - \
                                      data['first_funding_at_temp']) / pd.Timedelta(days=365)
```

```
In [39]: # checking for nulls
print(data['first_funding_at'][data['time_to_first_funding'].isnull()])
print(data['last_funding_at'][data['time_first_to_last_funding'].isnull()])
```

```
1030      0020-06-14
4514      0019-11-20
9863      0201-01-01
20287     0007-05-13
21784     0001-05-14
Name: first_funding_at, dtype: object
1030      2013-06-01
4514      2013-04-01
9863      0201-01-01
20287     2014-09-25
21784     0001-05-14
Name: last_funding_at, dtype: object
```


```
In [40]: # dropping these
data = data.dropna(subset=['time_to_first_funding'])
```

## OPTIONAL - Outliers

```
In [41]: # Funding total USD
print(np.percentile(data['funding_total_usd'], 95))

plt.hist(data['funding_total_usd'][data['funding_total_usd'] < 408206869399.0])
print(len(data[data['funding_total_usd'] < 408206869399.0]))
print(len(data))

408206869399.0
32338
34040
```




```
In [42]: np.percentile(data['venture'], 95)

plt.hist(data['venture'][data['venture'] < 40000000.0])
print(len(data[data['venture'] < 40000000.0]))
print(len(data[(data['venture'] < 40000000.0) &
                (data['funding_total_usd'] < 408206869399.0)]))
print(len(data))
```

32301

31990

34040



```
In [43]: ### THIS DOESN'T HELP THE MODELS, SO I COMMENTED THIS OUT
```

```
In [44]: # uncomment to remove outliers
```

```
#data = data[(data['venture'] < 40000000.0) &
#            (data['funding_total_usd'] < 408206869399.0)].reset_index(drop=True)
#len(data)
```

## Correlations

Based on the below, the strongest correlations occur between debt\_financing and funding\_total\_usd, round\_H and round\_G, followed by venture and all of the rounds of


## funding

```
In [45]: # with status
encoder = LabelEncoder()
heatmap_data = pd.concat([pd.Series(encoder.fit_transform(data['status'])),
                           data.drop('status', axis=1)], axis=1)

# without status
#heatmap_data = data.drop('status', axis=1)

h_corr = heatmap_data.corr()
fig, ax = plt.subplots(figsize=(20, 20))
mask = np.triu(np.ones_like(h_corr, dtype=bool))
cmap = sns.diverging_palette(130, 50, as_cmap=True)
cbar_kws = {'label': 'Correlation', 'orientation': 'horizontal',
            'pad': .1, 'extend': 'both'}
sns.heatmap(data=h_corr, mask=mask, ax=ax, annot=True, cbar_kws=cbar_kws, cmap=
ax.set_title('Heatmap of Correlation Between Attributes (Including Status)')
```

Out[45]: Text(0.5, 1.0, 'Heatmap of Correlation Between Attributes (Including Status)')



```
In [46]: #data_corr=data.drop(columns=['status'], axis=1).corr()
corr = h_corr.abs().stack().reset_index().sort_values(0, ascending=False)
corr['pairs'] = list(zip(corr.level_0, corr.level_1))
corr.set_index(['pairs'], inplace = True)
corr.drop(columns=['level_1', 'level_0'], inplace = True)
```

```
# cc for correlation coefficient
corr.columns = ['cc']
corr.drop_duplicates(inplace=True)

corr[(corr['cc'] > 0.7) & (corr['cc'] < 1)]
```

Out[46]:

cc

	cc
(round_H, round_G)	0.859849
(funding_rounds, time_first_to_last_funding)	0.793787

```
In [47]: # dropping temp columns
data = data.drop(columns=['founded_at', 'first_funding_at', 'last_funding_at',
                          'first_funding_at_temp', 'last_funding_at_temp',
                          'founded_at_temp'], axis=1).reset_index(drop=True)

# creating csv file to work from
data.to_csv('data/final_working_data.csv')
```

```
In [48]: # dropping columns with correlation coefficient greater than 0.7
data_uncorr = data.drop(columns=['round_H', 'time_first_to_last_funding'],
                          axis=1).reset_index(drop=True)

#renaming full dataset
data_full = data
```

## OPTIONAL - Binary Representation of Funding Rounds

```
In [49]: # data['had_round_A'] = [0 if x==0 else 1 for x in data['round_A']]
# data['had_round_B'] = [0 if x==0 else 1 for x in data['round_B']]
# data['had_round_C'] = [0 if x==0 else 1 for x in data['round_C']]
# data['had_round_D'] = [0 if x==0 else 1 for x in data['round_D']]
# data['had_round_E'] = [0 if x==0 else 1 for x in data['round_E']]
# data['had_round_F'] = [0 if x==0 else 1 for x in data['round_F']]
# data['had_round_G'] = [0 if x==0 else 1 for x in data['round_G']]
# data['had_venture'] = [0 if x==0 else 1 for x in data['venture']]
# data['had_seed'] = [0 if x==0 else 1 for x in data['seed']]
# data['had_eq_crowdfunding'] = [0 if x==0 else 1 for x in data['equity_crowdfu
# data['had_pd_crowdfunding'] = [0 if x==0 else 1 for x in data['product_crowdf
# data['had_angel'] = [0 if x==0 else 1 for x in data['angel']]
# data['had_grant'] = [0 if x==0 else 1 for x in data['grant']]
# data['had_pe'] = [0 if x==0 else 1 for x in data['private_equity']]
# data['had_convert'] = [0 if x==0 else 1 for x in data['convertible_note']]
```

## Functions

### train\_test\_preprocess

```
In [50]: '''
Function to perform train_test_split and necessary preprocessing / scaling
'''

def train_test_preprocess(X, y):
```



```

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)

# check that there are the same number of rows in X as values in y
assert X_train.shape[0] == y_train.shape[0]

# Categorizing features in preparation for scaling / encoding
X_train_cat_eng = X_train.select_dtypes(include=['int64']).reset_index(drop=True)
X_test_cat_eng = X_test.select_dtypes(include=['int64']).reset_index(drop=True)

X_train_cont = X_train.select_dtypes(exclude=['object', 'int64']).reset_index(drop=True)
X_test_cont = X_test.select_dtypes(exclude=['object', 'int64']).reset_index(drop=True)

cat_columns = ['market', 'region']
cat_train = X_train[cat_columns].reset_index(drop=True)
cat_test = X_test[cat_columns].reset_index(drop=True)

# Scale continuous variables using Min Max Scaler:
scaler = MinMaxScaler() # instantiate MinMaxScaler

## TRAIN
# Fit and transform X_train
X_train_cont_scaled = scaler.fit_transform(X_train_cont)
X_train_cont_scaled = pd.DataFrame(X_train_cont_scaled, columns=X_train_cont.columns)

# One hot encode categoricals
ohe = OneHotEncoder(handle_unknown = 'ignore')
encoded_train = ohe.fit_transform(cat_train).toarray()
X_train_cat = pd.DataFrame(encoded_train, columns=ohe.get_feature_names(cat_train.columns))

# Putting it all together:
X_train_processed = pd.concat([X_train_cat, X_train_cont, X_train_cat_eng], axis=1)
X_train_scaled = pd.concat([X_train_cat, X_train_cont_scaled, X_train_cat_eng], axis=1)

## TEST
# Scale continuous features
X_test_cont_scaled = scaler.transform(X_test_cont)
X_test_cont_scaled = pd.DataFrame(X_test_cont_scaled, columns=X_test_cont.columns)

# One hot encoding categoricals
encoded_test = ohe.transform(cat_test).toarray()
X_test_cat = pd.DataFrame(encoded_test, columns=ohe.get_feature_names(cat_test.columns))

# Putting it all together
X_test_scaled = pd.concat([X_test_cat, X_test_cont_scaled, X_test_cat_eng], axis=1)
X_test_processed = pd.concat([X_test_cat, X_test_cont, X_test_cat_eng], axis=1)

return X_train_processed, X_train_scaled, X_test_processed, X_test_scaled,

```

## print\_scores

```

In [51]: '''
Function to print relevant scoring metrics
'''

def print_scores(y_train, y_hat_train, y_test, y_hat_test, binary=True):
    if binary:
        print('Training Recall: ',
              recall_score(y_train, y_hat_train))
        print('Testing Recall: ',

```

```

        recall_score(y_test, y_hat_test))
print('\n')
print('Training F1: ',
      f1_score(y_train, y_hat_train))
print('Testing F1: ',
      f1_score(y_test, y_hat_test))
print('\n')
false_positive_rate, true_positive_rate, thresholds = \
roc_curve(y_test, y_hat_test)
roc_auc = auc(false_positive_rate, true_positive_rate)
print('ROC AUC: ', roc_auc)
print('PR AUC: ', average_precision_score(y_test, y_hat_test))
print('\n')

else:
    print('Training Recall (weighted avg): ',
          recall_score(y_train, y_hat_train, average='weighted'))
    print('Testing Recall (weighted avg): ',
          recall_score(y_test, y_hat_test, average='weighted'))
    print('\n')
    print('Training Recall (macro avg): ',
          recall_score(y_train, y_hat_train, average='macro'))
    print('Testing Recall (macro avg): ',
          recall_score(y_test, y_hat_test, average='macro'))
    print('\n')
    print('Training F1-Score (weighted avg): ',
          f1_score(y_train, y_hat_train, average='weighted'))
    print('Testing F1-Score (weighted avg): ',
          f1_score(y_test, y_hat_test, average='weighted'))
    print('\n')
    print('Training F1-Score (macro avg): ',
          f1_score(y_train, y_hat_train, average='macro'))
    print('Testing F1-Score (macro avg): ',
          f1_score(y_test, y_hat_test, average='macro'))
    print('\n')
    print('Testing Recall (failure class): ',
          recall_score(y_test, y_hat_test, average=None, labels=[1]))
    print('\n')

print('Training Accuracy: ', accuracy_score(y_train, y_hat_train))
print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))

```

## return\_scores

```

In [52]: '''
Function that stores relevant scoring metrics
'''

def return_scores(y_train, y_hat_train, y_test, y_hat_test):
    r_train = recall_score(y_train, y_hat_train)
    r_test = recall_score(y_test, y_hat_test)

    f1_train = f1_score(y_train, y_hat_train)
    f1_test = f1_score(y_test, y_hat_test)

    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_hat_test)
    roc_auc = auc(false_positive_rate, true_positive_rate)

    pr_auc = average_precision_score(y_test, y_hat_test)

```

```

ac_train = accuracy_score(y_train, y_hat_train)
ac_test = accuracy_score(y_test, y_hat_test)

return r_train, r_test, f1_train, f1_test, ac_train, ac_test, roc_auc, pr_a

```

## train\_test\_check

```

In [53]: '''
Function that checks new train & test splits for proper shape
'''

def train_test_check(X_train_processed, X_train_scaled, X_test_processed,
                     X_test_scaled, y_train, y_test):

    assert X_train_processed.shape[0] == y_train.shape[0]
    assert X_train_scaled.shape[0] == y_train.shape[0]

    assert X_test_processed.shape[0] == y_test.shape[0]
    assert X_test_scaled.shape[0] == y_test.shape[0]

    print("There are {} features in train set".format(len(X_train_processed.columns)))
    print("There are {} features in test set".format(len(X_test_processed.columns)))
    print('\n')

    print("There are {} features in train set (scaled)".format(len(X_train_scaled.columns)))
    print("There are {} features in test set (scaled)".format(len(X_test_scaled.columns)))
    print('\n')

    print(f"y_train is a Series with {y_train.shape[0]} values")
    print('\n')
    print("target breakdown: ", y_train.value_counts(normalize=True))

    display(X_train_processed.head())
    display(X_train_scaled.head())

```

## correlation\_check

```

In [54]: '''
Function that checks for excessive correlations across features
'''

def correlation_check(X_train_processed):

    df_corr=X_train_processed.corr()

    df = df_corr.abs().stack().reset_index().sort_values(0, ascending=False)
    df['pairs'] = list(zip(df.level_0, df.level_1))
    df.set_index(['pairs'], inplace = True)
    df.drop(columns=['level_1', 'level_0'], inplace = True)

    # cc for correlation coefficient
    df.columns = ['cc']
    df.drop_duplicates(inplace=True)

    display(df[(df.cc>.5) & (df.cc<1)])

```

# Acquired or Closed Subset

```
In [55]: data_ac = df[df['status'] != 'operating']
```

## EDA / Preprocessing

```
In [56]: data_ac.isnull().sum()
```

```
Out[56]: name                2
market              611
funding_total_usd    1
status             1315
country_code        835
state_code         2582
region             835
city               951
funding_rounds       1
founded_at         1720
founded_month       1725
founded_year        1725
first_funding_at     1
last_funding_at      1
seed                1
venture             1
equity_crowdfunding  1
undisclosed          1
convertible_note     1
debt_financing       1
angel               1
grant               1
private_equity       1
product_crowdfunding 1
round_A             1
round_B             1
round_C             1
round_D             1
round_E             1
round_F             1
round_G             1
round_H             1
dtype: int64
```

```
In [57]: len(data_ac)
```

```
Out[57]: 7610
```

```
In [58]: data_ac = data_ac.dropna(subset=['status', 'name'])
len(data_ac)
```

```
Out[58]: 6294
```

```
In [59]: data_ac.dtypes
```

```
Out[59]: name                object
market                object
funding_total_usd     object
status                object
country_code          object
state_code            object
region                object
city                  object
funding_rounds        float64
founded_at            object
founded_month         object
founded_year          float64
first_funding_at      object
last_funding_at       object
seed                  float64
venture                float64
equity_crowdfunding    float64
undisclosed            float64
convertible_note       float64
debt_financing         float64
angel                  float64
grant                  float64
private_equity         float64
product_crowdfunding   float64
round_A                float64
round_B                float64
round_C                float64
round_D                float64
round_E                float64
round_F                float64
round_G                float64
round_H                float64
dtype: object
```

```
In [60]: # converting to float
data_ac['funding_total_usd'] = [float(num) for num in data_ac['funding_total_usd']]
```

```
In [61]: data_ac['status'].value_counts(normalize=True)
```

```
Out[61]: acquired    0.58659
closed             0.41341
Name: status, dtype: float64
```


```
In [62]: # filling categoricals
data_ac = data_ac.fillna(value={'market': 'other', 'country_code': 'other',
                                'region': 'other', 'city': 'other',
                                'state_code': 'other'})
```

```
In [63]: # distribution of non NA years
plt.hist(data_ac['founded_year'][~data_ac['founded_year'].isnull()], bins=30)
plt.axvline(x=np.nanmean(data_ac['founded_year']), color='r')
plt.axvline(x=np.nanmedian(data_ac['founded_year']), color='y')

p1 = np.percentile(data_ac['founded_year'][~data_ac['founded_year'].isnull()],
print("1st percentile of founded_year: ", p1)

#1980 looks like a good cutoff point
print("Num Rows before 1980: ", len(data_ac[data_ac['founded_year'] < 1980]))
data_ac = data_ac[data_ac['founded_year'] >= 1980].reset_index(drop=True)
```

1st percentile of founded\_year: 1982.0  
 Num Rows before 1980: 35



```
In [64]: data_ac['founded_at'] = [datetime.strptime(day, '%Y-%m-%d') for day \
                                in data_ac['founded_at'][-data_ac['founded_at'].isnull()]]
data_ac['first_funding_at'] = [datetime.strptime(day, '%Y-%m-%d').date() for day \
                                in data_ac['first_funding_at']]
data_ac['last_funding_at'] = [datetime.strptime(day, '%Y-%m-%d').date() for day \
                                in data_ac['last_funding_at']]
data_ac['founded_month'] = [datetime.strptime(mth, "%Y-%m").month for mth in data_ac['founded_at']]
```

```
In [65]: # imputing median for day, then pulling founded year and month from that
imp_median = SimpleImputer(missing_values=np.nan, strategy='median')

data_ac['founded_at'] = imp_median.fit_transform(data_ac[['founded_at']])

data_ac['founded_at'] = pd.to_datetime(data_ac['founded_at'])
data_ac['founded_year'] = [day.year for day in data_ac['founded_at']]
data_ac['founded_month'] = [day.month for day in data_ac['founded_at']]
```

```
In [66]: # creating column that labels country as domestic or international
data_ac['international'] = [0 if country=='USA' else 1 for country in data_ac['country']]
```

```
In [67]: # creating temporary columns to aid in calculation of time to first funding

data_ac['founded_at_temp'] = [day.date() for day in data_ac['founded_at']]
data_ac['founded_at_temp'] = pd.to_datetime(data_ac['founded_at_temp'],
                                             format = '%Y-%m-%d')

data_ac['first_funding_at_temp'] = pd.to_datetime(data_ac['first_funding_at'],
                                                  format = '%Y-%m-%d',
                                                  errors='coerce')
data_ac['last_funding_at_temp'] = pd.to_datetime(data_ac['last_funding_at'],
                                                  format = '%Y-%m-%d',
                                                  errors='coerce')

data_ac['time_to_first_funding'] = (data_ac['first_funding_at_temp'] - \
                                   data_ac['founded_at_temp']) / pd.Timedelta('1d')

data_ac['time_first_to_last_funding'] = (data_ac['last_funding_at_temp'] - \
                                         data_ac['first_funding_at_temp']) / pd.Timedelta('1d')
```

```
In [68]: # checking for nulls
```

```
print(data_ac['first_funding_at'][data_ac['time_to_first_funding'].isnull()])
print(data_ac['last_funding_at'][data_ac['time_first_to_last_funding'].isnull()])
```

```
3697      0011-11-14
Name: first_funding_at, dtype: object
3697      2012-07-24
Name: last_funding_at, dtype: object
```

```
In [69]: data_ac = data_ac.dropna(subset=['time_to_first_funding']).reset_index(drop=True)
```

```
In [70]: # getting rid of extra spaces
```

```
data_ac['market'] = [x.strip() for x in data_ac['market']]
data_ac['country_code'] = [x.strip() for x in data_ac['country_code']]
data_ac['state_code'] = [x.strip() for x in data_ac['state_code']]
data_ac['region'] = [x.strip() for x in data_ac['region']]
data_ac['city'] = [x.strip() for x in data_ac['city']]
```

```
In [71]: # correlation check
```

```
# with status
```


```
encoder = LabelEncoder()
heatmap_data = pd.concat([pd.Series(encoder.fit_transform(data_ac['status'])),
                          data_ac.drop('status', axis=1)], axis=1)
```

```
# without status
```

```
#heatmap_data = data.drop('status', axis=1)
```

```
h_corr = heatmap_data.corr()
fig, ax = plt.subplots(figsize=(20, 20))
mask = np.triu(np.ones_like(h_corr, dtype=bool))
cmap = sns.diverging_palette(130, 50, as_cmap=True)
cbar_kws = {'label': 'Correlation', 'orientation': 'horizontal',
            'pad': .1, 'extend': 'both'}
sns.heatmap(data=h_corr, mask=mask, ax=ax, annot=True, cbar_kws=cbar_kws, cmap=cmap)
ax.set_title('Heatmap of Correlation Between Attributes (Including Status)')
```

```
Out[71]: Text(0.5, 1.0, 'Heatmap of Correlation Between Attributes (Including Status)')
```



```
In [72]: corr = h_corr.abs().stack().reset_index().sort_values(0, ascending=False)
corr['pairs'] = list(zip(corr.level_0, corr.level_1))
corr.set_index(['pairs'], inplace = True)
corr.drop(columns=['level_1', 'level_0'], inplace = True)

# cc for correlation coefficient
corr.columns = ['cc']
corr.drop_duplicates(inplace=True)

corr[(corr['cc'] > 0.7) & (corr['cc'] < 1)]
```

```
Out[72]:
```

	cc
(time_to_first_funding, founded_year)	0.819908
(funding_rounds, time_first_to_last_funding)	0.812019

```
In [73]: # reducing number of categories in market & region

print(sum(data_ac['market'].value_counts() >= 15)) # top ~50 markets
data_ac['market'] = data_ac['market'].map(lambda x: x if x >= 15 else 'Other')
print(data_ac['market'].value_counts())

print(sum(data_ac['region'].value_counts() >= 60)) # top 10-15 regions
```



```
data_ac['region'][data_ac['region'].map(data_ac['region'].value_counts()) < 60]  
print(data_ac['region'].value_counts())
```

```

48
other                1162
Software             569
Curated Web        347
Mobile              281
Enterprise Software 200
Biotechnology       189
Advertising         173
Games               172
E-Commerce          164
Social Media        136
Hardware + Software 100
Semiconductors      96
Security            92
Web Hosting         84
Clean Technology    78
Health Care         75
Finance             71
Analytics           59
Messaging           55
Search              54
News                47
Music               45
Education           43
Public Relations    42
Video               41
Travel              40
Networking           38
Photography         36
Social Network Media 34
Consulting          33
Health and Wellness 32
SaaS                29
Sports              25
Sales and Marketing 23
Web Development     23
Internet            23
Manufacturing       21
Cloud Computing     20
Android             20
iPhone              19
Fashion             19
Apps                18
Marketplaces        18
Facebook Applications 18
Automotive          17
Hospitality         17
Shopping            17
Real Estate         15
Name: market, dtype: int64
13
other                2128
SF Bay Area          1250
New York City        338
Boston               260
Los Angeles          186
London               146
Seattle              143
Washington, D.C.     85
Chicago              85

```

```

Austin            84
San Diego         77
Denver            75
Tel Aviv          73
Name: region, dtype: int64

```

```
In [74]: #data_ac['region'].value_counts()[data_ac['region'].value_counts() > 75]
```

```

In [75]: # dropping correlated columns for logistic regression, which can be
# sensitive to correlated features
data_final = data_ac.drop(columns=['founded_at', 'first_funding_at', 'last_funding_at',
                                   'first_funding_at_temp', 'last_funding_at_temp',
                                   'founded_at_temp', 'round_H', 'founded_year',
                                   'time_first_to_last_funding'],
                           axis=1)

# all inclusive
data_final2 = data_ac.drop(columns=['founded_at', 'first_funding_at', 'last_funding_at',
                                   'first_funding_at_temp', 'last_funding_at_temp',
                                   'founded_at_temp', 'round_H'],
                           axis=1)

```

```
In [76]: data_final.columns
```

```

Out[76]: Index(['name', 'market', 'funding_total_usd', 'status', 'country_code',
               'state_code', 'region', 'city', 'funding_rounds', 'founded_month',
               'seed', 'venture', 'equity_crowdfunding', 'undisclosed',
               'convertible_note', 'debt_financing', 'angel', 'grant',
               'private_equity', 'product_crowdfunding', 'round_A', 'round_B',
               'round_C', 'round_D', 'round_E', 'round_F', 'round_G', 'international',
               'time_to_first_funding'],
              dtype='object')

```

```
In [77]: data_final2.columns
```

```

Out[77]: Index(['name', 'market', 'funding_total_usd', 'status', 'country_code',
               'state_code', 'region', 'city', 'funding_rounds', 'founded_month',
               'founded_year', 'seed', 'venture', 'equity_crowdfunding', 'undisclosed',
               'convertible_note', 'debt_financing', 'angel', 'grant',
               'private_equity', 'product_crowdfunding', 'round_A', 'round_B',
               'round_C', 'round_D', 'round_E', 'round_F', 'round_G', 'international',
               'time_to_first_funding', 'time_first_to_last_funding'],
              dtype='object')

```

## X/Y Split

```

In [78]: X_ac = data_final2.drop(columns=['status', 'name', 'country_code',
                                   'state_code', 'city'], axis=1)

encoder = LabelEncoder()
y_ac = pd.Series(encoder.fit_transform(data_final2['status']))
# acquired is 0, closed is 1

y_ac.value_counts(normalize=True)

```

```

Out[78]: 0    0.59716
         1    0.40284
         dtype: float64

```

For the logistic regression model, we'll get rid of some extre outliers and exclude highly correlated features


```
In [79]: # Funding total USD
print(np.percentile(data_final['funding_total_usd'], 97.5))

plt.hist(data_final['funding_total_usd'][data_final['funding_total_usd'] < 80000000000.0])
print(len(data_final[data_final['funding_total_usd'] >= 80000000000.0]))
print(len(data_final))
```

807015451469.2499

147

4930



```
In [80]: print(np.percentile(data_final['venture'], 97.5))


plt.hist(data_final['venture'][data_final['venture'] < 70000000.0])
print(len(data_final[data_final['venture'] > 70000000.0]))
print(len(data_final[(data_final['venture'] > 70000000.0) &
                        (data_final['funding_total_usd'] > 80000000000.0)]))
print(len(data_final))
```

69561028.64999995

119

96

4930



```
In [81]: data_lr = data_final[(data_final['venture'] < 70000000.0) &
                             (data_final['funding_total_usd'] < 80000000000.0)].reset_
```

```
In [82]: correlation_check(data_lr)
```

cc	
pairs	
(venture, funding_total_usd)	0.920258
(round_B, venture)	0.612658
(venture, round_C)	0.601337
(funding_total_usd, round_C)	0.562161
(funding_total_usd, round_B)	0.557115
(venture, funding_rounds)	0.537058
(funding_total_usd, funding_rounds)	0.509066

```
In [83]: x_ac_lr = data_lr.drop(columns=['status', 'name', 'country_code',
                                         'state_code', 'city', 'funding_total_usd'],
                               axis=1)

encoder = LabelEncoder()
y_ac_lr = pd.Series(encoder.fit_transform(data_lr['status']))
# acquired is 0, closed is 1

y_ac_lr.value_counts(normalize=True)
```

```
Out[83]: 0    0.588742
         1    0.411258
         dtype: float64
```

## Baseline Model

```
In [84]: X_train_processed, X_train_scaled, X_test_processed, \
         X_test_scaled, y_train, y_test = train_test_preprocess(X_ac_lr, y_ac_lr)
```

```
In [85]: train_test_check(X_train_processed, X_train_scaled, X_test_processed,
                          X_test_scaled, y_train, y_test)
```

There are 82 features in train set  
There are 82 features in test set

There are 82 features in train set (scaled)  
There are 82 features in test set (scaled)

y\_train is a Series with 3570 values

```
target breakdown: 0    0.581793
                  1    0.418207
                  dtype: float64
```

	market_Advertising	market_Analytics	market_Android	market_Apps	market_Automotive	n
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	0.0	

5 rows × 82 columns

	market_Advertising	market_Analytics	market_Android	market_Apps	market_Automotive	n
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	0.0	

5 rows × 82 columns

```
In [86]: baseline_model = LogisticRegression(random_state=42)
baseline_model.fit(X_train_scaled, y_train)
```

```
Out[86]: LogisticRegression(random_state=42)
```

```
In [87]: y_hat_train = baseline_model.predict(X_train_scaled)
y_hat_test = baseline_model.predict(X_test_scaled)

print_scores(y_train, y_hat_train, y_test, y_hat_test)
```

```
Training Recall: 0.6262558606831882
Testing Recall: 0.6
```

```
Training F1: 0.6572934973637962
Testing F1: 0.6058631921824104
```

```
ROC AUC: 0.6780991735537191
PR AUC: 0.5232765477926554
```

```
Training Accuracy: 0.726890756302521
Testing Accuracy: 0.6952141057934509
```

## Baseline Analysis

- 60% recall is not a great start, but would like to see if hypertuning can improve this
- 70% accuracy also leaves something to be desired

## Grid Search on Baseline\*\*

```
In [88]: param_grid = {'penalty':['l2', None],
                      'solver':['lbfgs', 'sag'],
                      'C': [1.0, 1e12],
                      'class_weight': [None, 'balanced']}
                      }
```

```
In [89]: grid_logreg = GridSearchCV(baseline_model, param_grid, cv = 5,
                                   scoring='recall') # macro or weighted
grid_logreg.fit(X_train_scaled, y_train)
grid_logreg.best_params_
```

```
Out[89]: {'C': 1000000000000.0,
          'class_weight': 'balanced',
          'penalty': 'l2',
          'solver': 'lbfgs'}
```

```
In [90]: y_preds_grid_lr_train = grid_logreg.predict(X_train_scaled)
y_preds_grid_lr = grid_logreg.predict(X_test_scaled)

print_scores(y_train, y_preds_grid_lr_train, y_test, y_preds_grid_lr)
```

```
Training Recall:  0.769591426657736
Testing Recall:  0.7569892473118279
```

```
Training F1:  0.6993304930006086
Testing F1:  0.6597938144329897
```

```
ROC AUC:  0.706318315115969
PR AUC:  0.5375031955182031
```

```
Training Accuracy:  0.7232492997198879
Testing Accuracy:  0.6952141057934509
```

```
In [91]: ## Top coefficients
lr = LogisticRegression(random_state=42, class_weight='balanced',
                        solver='lbfgs', C=1e12, penalty='l2')
lr.fit(X_train_scaled, y_train)

coef_df = pd.DataFrame(lr.coef_, columns=X_train_scaled.columns).transpose()
coef_df.to_csv('coef_logreg.csv')

coef_df.columns=['coef']
coef_df['coef_abs'] = abs(coef_df['coef'])

coef_df.sort_values(by='coef_abs', ascending=False)[:20]
```

Out [91]:

	coef	coef_abs
<b>time_to_first_funding</b>	-13.318211	13.318211
<b>private_equity</b>	-4.029339	4.029339
<b>funding_rounds</b>	-3.113184	3.113184
<b>undisclosed</b>	-3.080005	3.080005
<b>seed</b>	-2.994236	2.994236
<b>convertible_note</b>	2.754146	2.754146
<b>venture</b>	-2.022008	2.022008
<b>market_Clean Technology</b>	2.021825	2.021825
<b>round_A</b>	-1.502466	1.502466
<b>market_Consulting</b>	1.492824	1.492824
<b>market_Analytics</b>	-1.264251	1.264251
<b>debt_financing</b>	-1.262151	1.262151
<b>market_Sports</b>	1.223371	1.223371
<b>market_Biotechnology</b>	1.194763	1.194763
<b>region_other</b>	1.120454	1.120454
<b>region_Denver</b>	0.919126	0.919126
<b>market_Hardware + Software</b>	0.860133	0.860133
<b>market_Apps</b>	-0.849601	0.849601
<b>round_F</b>	0.849171	0.849171
<b>grant</b>	-0.826037	0.826037

```
In [92]: X_train_scaled['private_equity'].value_counts()
# vast majority of observations are 0, so coefficient probably isn't
# as important as its ranking implies
```



```
Out[92]: 0.000000    3529
0.026667      2
0.932000      1
0.006667      1
0.800000      1
0.133333      1
0.126667      1
0.933333      1
0.011106      1
0.840000      1
0.046667      1
0.033333      1
0.123335      1
0.022267      1
0.160000      1
0.533333      1
0.960000      1
0.000173      1
1.000000      1
0.408116      1
0.148482      1
0.529627      1
0.042667      1
0.248854      1
0.186667      1
0.252775      1
0.019813      1
0.053333      1
0.074248      1
0.060000      1
0.154667      1
0.240000      1
0.087976      1
0.066667      1
0.005051      1
0.081411      1
0.244853      1
0.624849      1
0.724281      1
0.006000      1
0.120000      1
Name: private_equity, dtype: int64
```

## Custom Pipeline

```
In [93]: scores = pd.DataFrame(columns = ['recall_train', 'recall_test', 'f1_train',
                                         'f1_test', 'accuracy_train', 'accuracy_test', \
                                         'roc_auc', 'pr_auc', 'params'])
```

```
In [94]: # creates a data frame with various scores for each model
def customPipe(model, model_name, X, y):
    X_train_processed, X_train_scaled, X_test_processed, X_test_scaled, \
    y_train, y_test = train_test_preprocess(X, y)

    if ('lr' in model_name) | ('knn' in model_name):
        model.fit(X_train_scaled, y_train)
        y_hat_train = model.predict(X_train_scaled)
        y_hat_test = model.predict(X_test_scaled)
```

```

else:
    model.fit(X_train_processed, y_train)
    y_hat_train = model.predict(X_train_processed)
    y_hat_test = model.predict(X_test_processed)

    r_train, r_test, f1_train, f1_test, ac_train, ac_test, roc_auc, pr_auc = \
    return_scores(y_train, y_hat_train, y_test, y_hat_test)

    score_list = []
    score_list.extend((r_train, r_test, f1_train, f1_test,
                       ac_train, ac_test, roc_auc, pr_auc, str(model)))

    scores.loc[model_name] = score_list
return scores

```

```

In [95]: # model inputs
lr = LogisticRegression(random_state=42, class_weight='balanced',
                        solver='lbfgs', C=1e12, penalty='l2')
rf = RandomForestClassifier(random_state=42)
dtc = DecisionTreeClassifier(random_state=42)
ext = ExtraTreesClassifier(random_state=42)
xgb = XGBClassifier(random_state=42)
knn = KNeighborsClassifier()

```

```

In [96]: # running the function
customPipe(lr, 'lr', X_ac_lr, y_ac_lr)
customPipe(rf, 'rf', X_ac, y_ac)
customPipe(dtc, 'dtc', X_ac, y_ac)
customPipe(ext, 'ext', X_ac, y_ac)
customPipe(xgb, 'xgb', X_ac, y_ac)
customPipe(knn, 'knn', X_ac, y_ac)

```

```

Out[96]:

```

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc_auc	
lr	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.706318	0.5
rf	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.712164	0.5
dtc	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.658468	0.5
ext	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.690884	0.5
xgb	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.706877	0.5
knn	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.644163	0.5

Looking at recall and accuracy, the Random Forest and XGBoost models seem to yield the best scores. Will hyper tune each of these models.

```

In [97]: # storing dataframe for easy access
scores_base = scores # models pre hyper parameter tuning

```

## Other Models - Hypertuning

```
In [98]: X_train_processed, X_train_scaled, X_test_processed, \
X_test_scaled, y_train, y_test = train_test_preprocess(X_ac, y_ac)
```

```
In [99]: train_test_check(X_train_processed, X_train_scaled, X_test_processed,
X_test_scaled, y_train, y_test)
```

There are 85 features in train set  
There are 85 features in test set

There are 85 features in train set (scaled)  
There are 85 features in test set (scaled)

y\_train is a Series with 3697 values

target breakdown: 0 0.597782  
1 0.402218  
dtype: float64

	market_Advertising	market_Analytics	market_Android	market_Apps	market_Automotive	n
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	0.0	

5 rows × 85 columns

	market_Advertising	market_Analytics	market_Android	market_Apps	market_Automotive	n
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	0.0	

5 rows × 85 columns

## Random Forest

```
In [100]: rf = RandomForestClassifier(random_state=42)

# Initial search
rf_param_grid = {
    'criterion': ['gini', 'entropy', 'log_loss'],
    'max_depth': [8, 12, 20],
    'min_samples_leaf': [5, 10],
    'class_weight': [None, 'balanced']
}
```

```

}

## results from initial search
# {'class_weight': 'balanced',
#  'criterion': 'gini',
#  'max_depth': 20,
#  'min_samples_leaf': 5}

# Fine tuning
rf_param_grid2 = {
    'criterion': ['gini'],
    'max_depth': [12, 20, 25],
    'min_samples_leaf': [10, 15, 20],
    'class_weight': ['balanced'],
    'n_estimators': [100, 200]
}

grid_rfc = GridSearchCV(rf, rf_param_grid2, cv = 5, scoring='recall')
grid_rfc.fit(X_train_processed, y_train)

grid_rfc.best_params_

```

Out[100]: {'class\_weight': 'balanced',  
 'criterion': 'gini',  
 'max\_depth': 12,  
 'min\_samples\_leaf': 20,  
 'n\_estimators': 100}

```

In [101... y_preds_grid_rfc_train = grid_rfc.predict(X_train_processed)
y_preds_grid_rfc = grid_rfc.predict(X_test_processed)

print_scores(y_train, y_preds_grid_rfc_train, y_test, y_preds_grid_rfc)

```

Training Recall: 0.7222595830531271  
 Testing Recall: 0.6973947895791583

Training F1: 0.7049556941253693  
 Testing F1: 0.6850393700787402

ROC AUC: 0.7335747789857645  
 PR AUC: 0.5918918112508587

Training Accuracy: 0.7568298620503111  
 Testing Accuracy: 0.740470397404704

```

In [102... rf_best = RandomForestClassifier(random_state = 42, class_weight='balanced',
                                         criterion='gini', max_depth=12,
                                         min_samples_leaf=20, n_estimators=100)
customPipe(rf_best, 'rf_best', X_ac, y_ac)

```

Out[102]:

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc_auc
<b>lr</b>	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.706318
<b>rf</b>	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.712164
<b>dtc</b>	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.658468
<b>ext</b>	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.690884
<b>xgb</b>	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.706877
<b>knn</b>	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.644163
<b>rf_best</b>	0.722260	0.697395	0.704956	0.685039	0.756830	0.740470	0.733575

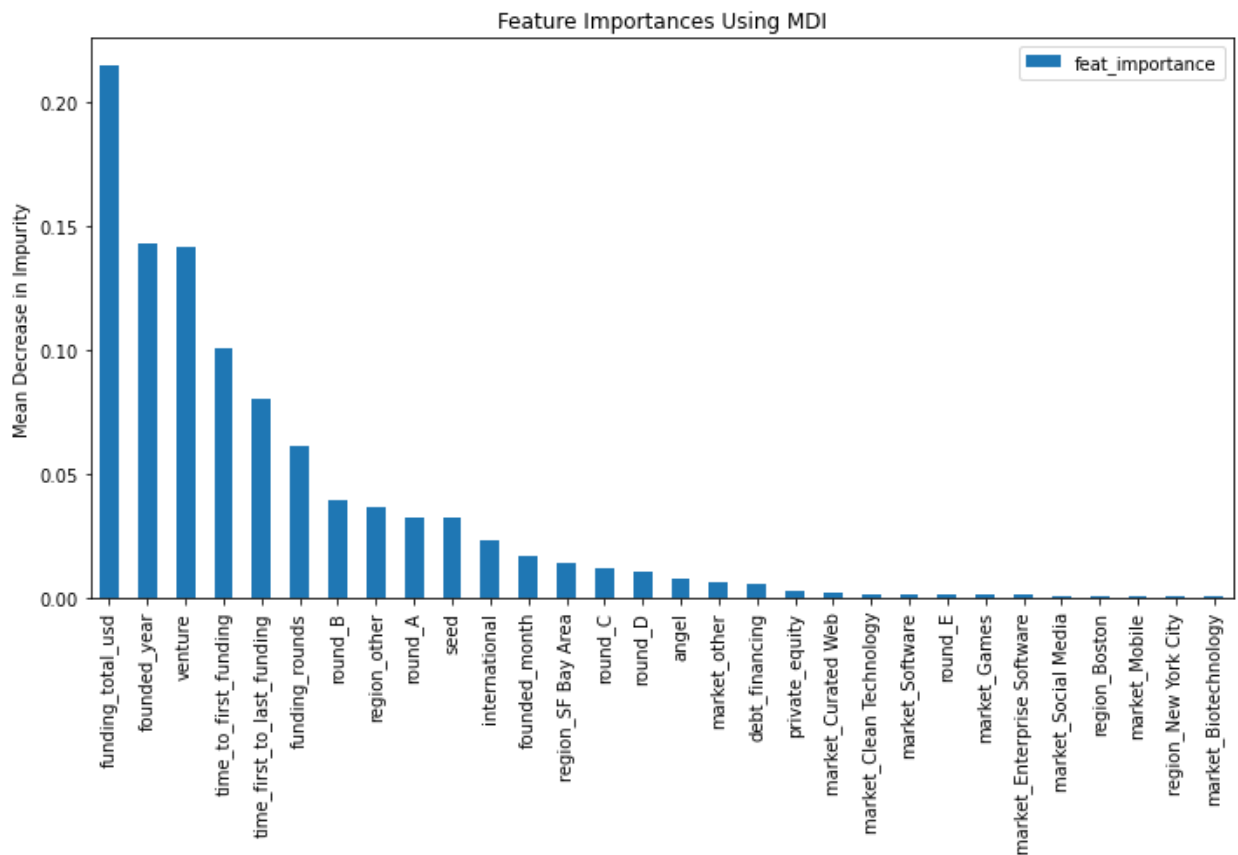
## Feature Importance

```
In [103... feats = {} # a dict to hold feature_name: feature_importance
for feature, importance in zip(X_train_processed.columns, rf_best.feature_importances_):
    feats[feature] = importance #add the name/value pair

feats_rf = pd.DataFrame(feats.items())

feats_rf.columns = ['col', 'feat_importance']
feats_rf = feats_rf.sort_values(by=['feat_importance'], ascending=False)
feats_rf_20 = feats_rf[:20]
feats_rf_30 = feats_rf[:30]
```

```
In [104... fig, ax = plt.subplots(figsize = (10,7))
feats_rf_30.plot.bar(ax=ax) # yerr=std,
ax.set_title("Feature Importances Using MDI")
ax.set_ylabel("Mean Decrease in Impurity")
ax.set_xticklabels(feats_rf_30['col'])
fig.tight_layout()
```



```
In [105]: list(feats_rf['col'][:10])
```

```
Out[105]: ['funding_total_usd',
            'founded_year',
            'venture',
            'time_to_first_funding',
            'time_first_to_last_funding',
            'funding_rounds',
            'round_B',
            'region_other',
            'round_A',
            'seed']
```

## XG Boost\*

```
In [106]: # 45+ runtime
xgb = XGBClassifier(random_state=42)

# initial grid search
xgb_param_grid = {
    'learning_rate': [0.1, 0.2],
    'max_depth': [4, 6, 8],
    'min_child_weight': [3, 5, 7],
    'subsample': [0.5, 0.7],
    'scale_pos_weight': [1.5, 2]
}

## result from initial grid search:
# {'learning_rate': 0.1,
#  'max_depth': 4,
#  'min_child_weight': 7,
```

```
# 'scale_pos_weight': 2,
# 'subsample': 0.7}

# fine tuning
xgb_param_grid2 = {
    'learning_rate': [0.1,0.2],
    'max_depth': [1,2,4],
    'min_child_weight': [3,7,10],
    'subsample': [0.7],
    'scale_pos_weight': [2]
}

grid_xgb = GridSearchCV(xgb, xgb_param_grid2, cv = 5, scoring='recall')
grid_xgb.fit(X_train_processed, y_train)

grid_xgb.best_params_
```

```
Out[106]: {'learning_rate': 0.2,
           'max_depth': 1,
           'min_child_weight': 7,
           'scale_pos_weight': 2,
           'subsample': 0.7}
```

```
In [107... xgb_best = XGBClassifier(random_state=42, max_depth=1, min_child_weight=7,
                           subsample=0.7, learning_rate=0.2, scale_pos_weight=2)

customPipe(xgb_best, 'xgb_best', X_ac, y_ac)
```

```
Out[107]:
```

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc_auc
<b>lr</b>	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.7063
<b>rf</b>	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.7121
<b>dtc</b>	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.6584
<b>ext</b>	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.6908
<b>xgb</b>	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.7068
<b>knn</b>	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.6441
<b>rf_best</b>	0.722260	0.697395	0.704956	0.685039	0.756830	0.740470	0.7335
<b>xgb_best</b>	0.803631	0.799599	0.706682	0.696943	0.731674	0.718573	0.7315

## Feature Importance

```
In [108... xgb_best.fit(X_train_processed, y_train)
```

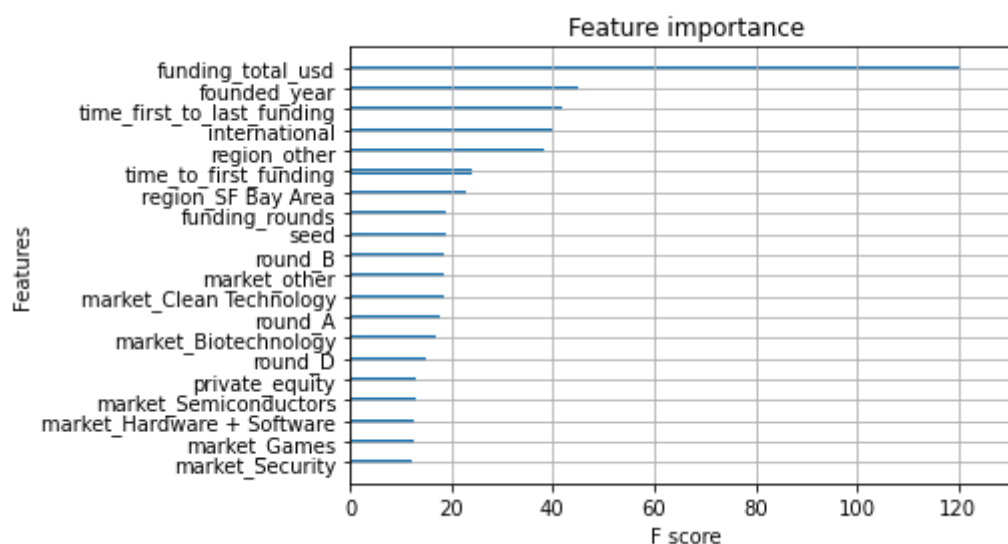
```
Out[108]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.2, max_delta_step=0, max_depth=1,
                        min_child_weight=7, missing=nan, monotone_constraints='()',
                        n_estimators=100, n_jobs=0, num_parallel_tree=1, random_state=4
                        2,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=2, subsample=0.7,
                        tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [109]: f = 'gain' # importance type
feat_imp = xgb_best.get_booster().get_score(importance_type= f)
feats_xg = pd.DataFrame(sorted(feat_imp.items(), key=lambda item: item[1],
                              reverse=True))
feats_xg.columns = ['col', 'feat_importance']
feats_xg[:15]
```

```
Out[109]:
```

	col	feat_importance
0	funding_total_usd	120.335433
1	founded_year	45.184254
2	time_first_to_last_funding	41.682578
3	international	39.813609
4	region_other	38.373383
5	time_to_first_funding	24.078237
6	region_SF Bay Area	22.668856
7	funding_rounds	18.984874
8	seed	18.831061
9	round_B	18.659117
10	market_other	18.592387
11	market_Clean Technology	18.477624
12	round_A	17.832843
13	market_Biotechnology	16.774651
14	round_D	15.123461

```
In [110]: from xgboost import plot_importance
plot_importance(xgb_best, max_num_features=20, importance_type=f,
               show_values=False)
plt.show()
```



```
In [111]: feats_xg[:15]
```



```

sns.set_style("whitegrid")
sns.set_palette('Blues')


g = sns.barplot(x='feat_importance', y='col', data=feats_xg[:15], color='pink')
# g.set_xticklabels(["U.S.", "International"])
g.set(xlabel='Feature Importance Score',
      ylabel='Feature',
      title="XG Boost Feature Importance")

```

```

Out[111]: [Text(0.5, 0, 'Feature Importance Score'),
Text(0, 0.5, 'Feature'),
Text(0.5, 1.0, 'XG Boost Feature Importance')]

```



## ExtraTrees

```

In [112... ext_param_grid = {'criterion':['entropy','gini'],
                        'max_depth':[15,20,25],
                        'min_samples_leaf': [1,5],
                        'class_weight': ['balanced'],
                        'max_features': ['auto']}

ext = ExtraTreesClassifier(random_state=42)

grid_ext = GridSearchCV(ext, ext_param_grid, cv = 5, scoring='recall')
grid_ext.fit(X_train_processed, y_train)

grid_ext.best_params_

```

```

Out[112]: {'class_weight': 'balanced',
'criterion': 'gini',
'max_depth': 15,
'max_features': 'auto',
'min_samples_leaf': 5}

```

```

In [113... ext_grid = ExtraTreesClassifier(random_state=42, class_weight='balanced',
                                         criterion='gini', max_depth= 15,
                                         max_features='auto', min_samples_leaf= 5)
customPipe(ext_grid, 'ext_best', X_ac, y_ac)

```

Out[113]:

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc_auc
<b>lr</b>	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.7063
<b>rf</b>	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.71210
<b>dtc</b>	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.65840
<b>ext</b>	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.69080
<b>xgb</b>	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.70680
<b>knn</b>	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.64410
<b>rf_best</b>	0.722260	0.697395	0.704956	0.685039	0.756830	0.740470	0.73350
<b>xgb_best</b>	0.803631	0.799599	0.706682	0.696943	0.731674	0.718573	0.73150
<b>ext_best</b>	0.759247	0.705411	0.711854	0.663525	0.752773	0.710462	0.70960

## Decision Tree\*\*\*

```
In [114]: dtc_param_grid = {'criterion':['gini','entropy'],
                             'max_depth':[5,10,15],
                             'min_samples_leaf': [15,20,25],
                             'class_weight': [None, 'balanced']}

dtc = DecisionTreeClassifier(random_state=42)

grid_dtc = GridSearchCV(dtc, dtc_param_grid, cv = 5, scoring='recall')
grid_dtc.fit(X_train_processed, y_train)

grid_dtc.best_params_
```

```
Out[114]: {'class_weight': 'balanced',
           'criterion': 'gini',
           'max_depth': 5,
           'min_samples_leaf': 15}
```

```
In [115]: dtc_best = DecisionTreeClassifier(random_state=42, class_weight='balanced',
                                             criterion='gini', max_depth= 5,
                                             min_samples_leaf= 15)

dtc_best.fit(X_train_processed, y_train)
customPipe(dtc_best, 'dtc_best', X_ac, y_ac)
```

Out[115]:

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc_auc
<b>lr</b>	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.7063
<b>rf</b>	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.71210
<b>dtc</b>	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.65840
<b>ext</b>	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.69080
<b>xgb</b>	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.70680
<b>knn</b>	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.64410
<b>rf_best</b>	0.722260	0.697395	0.704956	0.685039	0.756830	0.740470	0.73350
<b>xgb_best</b>	0.803631	0.799599	0.706682	0.696943	0.731674	0.718573	0.73150
<b>ext_best</b>	0.759247	0.705411	0.711854	0.663525	0.752773	0.710462	0.70960
<b>dtc_best</b>	0.806994	0.787575	0.710900	0.687063	0.736002	0.709651	0.72210


## Feature Importance

```
In [116... feats = {} # a dict to hold feature_name: feature_importance
for feature, importance in zip(X_train_processed.columns, dtc_best.feature_importances_):
    feats[feature] = importance #add the name/value pair

feats_dtc = pd.DataFrame(feats.items())

feats_dtc.columns = ['col', 'feat_importance']
feats_dtc = feats_dtc.sort_values(by=['feat_importance'], ascending=False)
feats_dtc_30 = feats_dtc[:30]
```

```
In [117... fig, ax = plt.subplots(figsize = (10,7))
feats_dtc_30.plot.bar(ax=ax) # yerr=std,
ax.set_title("Feature importances using MDI")
ax.set_ylabel("Mean decrease in impurity")
ax.set_xticklabels(feats_dtc_30['col'])
fig.tight_layout()
```



## Final Model

- XG Boost is the best model, with ~80% recall, 72% accuracy and 73% AUC
- Now that we have narrowed down the best models, we will run them with reduced features & SMOTE to see if that generates any improvement

```
In [118... X_train_processed, X_train_scaled, X_test_processed, \
X_test_scaled, y_train, y_test = train_test_preprocess(X_ac, y_ac)
```

## Feature Reduction

```
In [119... # top 20 features
X_train_xg = X_train_processed[list(feats_xg['col'][:20])]
X_train_dtc = X_train_processed[list(feats_dtc['col'][:20])]

X_test_xg = X_test_processed[list(feats_xg['col'][:20])]
X_test_dtc = X_test_processed[list(feats_dtc['col'][:20])]
```

```
In [120... # XG Boost
final_model = XGBClassifier(random_state=42, max_depth=1, min_child_weight=7,
                             subsample=0.7, learning_rate=0.2, scale_pos_weight=

final_model.fit(X_train_xg, y_train)

y_hat_train = final_model.predict(X_train_xg)
y_hat_test = final_model.predict(X_test_xg)
```

```
print_scores(y_train, y_hat_train, y_test, y_hat_test)
```

Training Recall: 0.8049764626765299

Testing Recall: 0.7875751503006012

Training F1: 0.7105966162065894

Testing F1: 0.6906854130052724

ROC AUC: 0.7262126432701916

PR AUC: 0.570346385974146

Training Accuracy: 0.7362726535028401

Testing Accuracy: 0.7145174371451744

```
In [121... # Decision Trees - 2nd place
final_model = dtc_best
final_model.fit(X_train_dtc, y_train)

y_hat_train = final_model.predict(X_train_dtc)
y_hat_test = final_model.predict(X_test_dtc)

print_scores(y_train, y_hat_train, y_test, y_hat_test)
```

Training Recall: 0.8069939475453934

Testing Recall: 0.7875751503006012

Training F1: 0.7109004739336493

Testing F1: 0.6870629370629371

ROC AUC: 0.7221254498097012

PR AUC: 0.565840551507965

Training Accuracy: 0.7360021639166892

Testing Accuracy: 0.7096512570965126

Feature reduction doesn't improve the models.

## SMOTE

```
In [122... #Using SMOTE to further reduce class imbalance
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_sample(X_train_processed, y_train)
```

```
In [123... # addressing class imbalance within the hyperparameters results in higher recall
final_model_smote = XGBClassifier(random_state=42, max_depth=1, min_child_weight=1,
                                   subsample=0.7, learning_rate=0.2)

final_model_smote.fit(X_train_resampled, y_train_resampled)

y_hat_train = final_model_smote.predict(X_train_resampled)
y_hat_test = final_model_smote.predict(X_test_processed)
```

```
print_scores(y_train_resampled, y_hat_train, y_test, y_hat_test)
```

Training Recall: 0.7773755656108597

Testing Recall: 0.7394789579158316

Training F1: 0.7712682379349045

Testing F1: 0.7001897533206831

ROC AUC: 0.7430364816827113

PR AUC: 0.5970874784794594

Training Accuracy: 0.7694570135746607

Testing Accuracy: 0.7437145174371452

## Model Selection

```
In [124]: final_model = XGBClassifier(random_state=42, max_depth=1, min_child_weight=7,
                                     subsample=0.7, learning_rate=0.2, scale_pos_weight=
                                     customPipe(final_model, 'final_model', X_ac, y_ac))
```

```
Out[124]:
```

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc
<b>lr</b>	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.70
<b>rf</b>	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.71
<b>dtc</b>	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.658
<b>ext</b>	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.690
<b>xgb</b>	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.700
<b>knn</b>	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.64
<b>rf_best</b>	0.722260	0.697395	0.704956	0.685039	0.756830	0.740470	0.73
<b>xgb_best</b>	0.803631	0.799599	0.706682	0.696943	0.731674	0.718573	0.73
<b>ext_best</b>	0.759247	0.705411	0.711854	0.663525	0.752773	0.710462	0.709
<b>dtc_best</b>	0.806994	0.787575	0.710900	0.687063	0.736002	0.709651	0.72
<b>final_model</b>	0.803631	0.799599	0.706682	0.696943	0.731674	0.718573	0.73

```
In [125]: # storing dataframe for easy access
scores_final = scores
scores_final.to_csv('final_scores.csv')
scores_final.sort_values(by=['recall_test'], ascending=False)[1:]
```

Out[125]:

	recall_train	recall_test	f1_train	f1_test	accuracy_train	accuracy_test	roc
<b>final_model</b>	0.803631	0.799599	0.706682	0.696943	0.731674	0.718573	0.73
<b>dtc_best</b>	0.806994	0.787575	0.710900	0.687063	0.736002	0.709651	0.72
<b>lr</b>	0.769591	0.756989	0.699330	0.659794	0.723249	0.695214	0.70
<b>ext_best</b>	0.759247	0.705411	0.711854	0.663525	0.752773	0.710462	0.70
<b>rf_best</b>	0.722260	0.697395	0.704956	0.685039	0.756830	0.740470	0.73
<b>xgb</b>	0.933423	0.607214	0.927188	0.641949	0.941033	0.725872	0.70
<b>rf</b>	1.000000	0.589178	1.000000	0.643326	1.000000	0.735604	0.71
<b>dtc</b>	1.000000	0.577154	1.000000	0.588957	1.000000	0.673966	0.65
<b>knn</b>	0.709482	0.577154	0.714528	0.576577	0.771977	0.656934	0.64
<b>ext</b>	1.000000	0.571142	1.000000	0.617551	1.000000	0.713706	0.69

In [126...

```

# visualizing predictions

y_hat_test = final_model.predict(X_test_processed)

conf_matrix = confusion_matrix(y_test, y_hat_test)

fig, ax = plt.subplots(figsize=(5,5))


ax = sns.heatmap(conf_matrix, annot=True, cmap='Greens', fmt='d')

ax.set_title('Final Model Confusion Matrix', fontsize=16);
ax.set_xlabel('Predicted Values', fontsize=13)
ax.set_ylabel('Actual Values ', fontsize=13);

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['Succeeded', 'Failed'])
ax.yaxis.set_ticklabels(['Succeeded', 'Failed'])

## Display the visualization of the Confusion Matrix.
plt.show()

```




## Final Analysis & Visualizations

```
In [127... ## Model Comparison
scores_viz = scores_final.sort_values(by=['recall_test'], ascending=False)[1:].

fig, ax = plt.subplots(figsize=(10, 5))
sns.stripplot(x="index", y="recall_test", data=scores_viz, size=10)
plt.xticks(rotation = 45)
ax.set_xlabel("Model", fontsize=10)
ax.set_ylabel("Recall", fontsize=10)
fig.suptitle("Model Performance", fontsize=15)
```

```
Out[127]: Text(0.5, 0.98, 'Model Performance')
```





In [128... *## Visualizing Important Features*

```
In [129... df = data_final2
df['int_category'] = ['international' if x==1 else 'U.S.' for x in df['international']
df['one_funding_round'] = ['one' if x==1 else 'multiple' for x in df['funding_rounds']]

df_box = df[(df['venture'] < 40000000.0)]

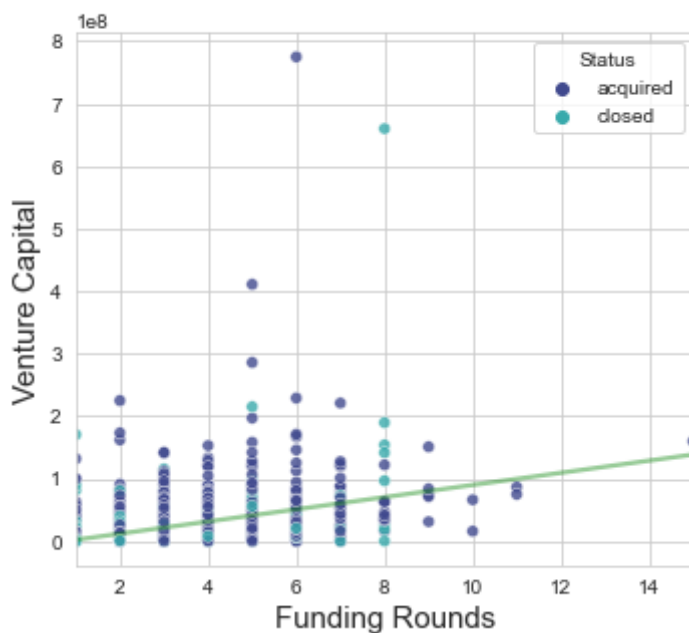
cmap = sns.diverging_palette(130, 50, as_cmap=True)
```

```
In [130... # setting up the figure with a single subplot
fig = plt.figure(figsize=(5, 5))
fig, ax = plt.subplots(figsize=(5, 5))
# scatterplot using seaborn
plot = sns.scatterplot(x='funding_rounds', y='venture', data=df,
hue=df.status, legend='full', alpha = 0.8, palette='mako')
# adding regression line using seaborn regplot
sns.regplot(data=df, x='funding_rounds', y='venture', scatter=False,
ax=ax, ci=False, color='g', line_kws={'alpha':0.4})

# updating figure title, adding labels for x- and y-axis
fig.suptitle("Venture vs Funding Rounds", fontsize=18)
ax.set_xlabel("Funding Rounds", fontsize=15)
ax.set_ylabel("Venture Capital", fontsize=15)
# setting legend title
ax.get_legend().set_title("Status")
# getting everything to fit nicely on the plot
plt.tight_layout()
```

<Figure size 360x360 with 0 Axes>

## Venture vs Funding Rounds



```
In [131]: fig = plt.figure(figsize=(5, 5))
fig, ax = plt.subplots(figsize=(5, 5))


df_fund = df[(df['funding_total_usd'] < 1000000) &
              (df['time_to_first_funding'] > -10) &
              (df['funding_total_usd'] > 0)]

# plot = sns.scatterplot(x='time_to_first_funding', y='funding_total_usd',
#                         data=df_fund, hue=df.status, legend='full',
#                         alpha = 0.8, palette='mako')

# histograms - funding used for successes & failures
df_success = df[df['status']=='acquired']
df_failure = df[df['status']=='closed']

plot = sns.histplot(x='funding_total_usd',
                    data=df_success[df_success['funding_total_usd'] < 750000],
                    bins=20)
```


<Figure size 360x360 with 0 Axes>



```
In [132]: df_box = df[(df['venture'] < 40000000.0)]
#df_box = df[(df['funding_total_usd'] < 1000000.0)]

plt.figure(figsize=(10,5))
sns.boxplot(x=df_box['venture'], y=df_box['status'])
```

Out[132]: <AxesSubplot:xlabel='venture', ylabel='status'>




```
In [133]: plt.figure(figsize=(10,5))

sns.boxplot(x=df_box['venture'], y=df_box['int_category'])

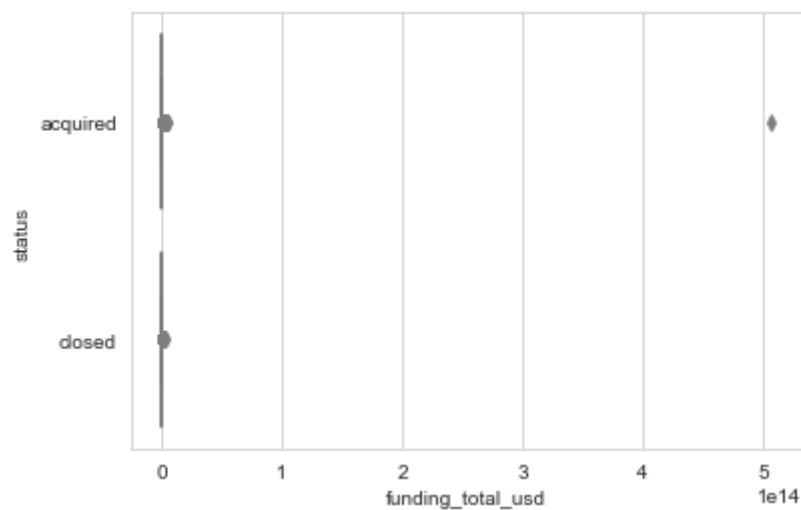
df_box['int_category'].value_counts()
```

Out[133]: U.S. 3143  
international 1443  
Name: int\_category, dtype: int64



```
In [134]: sns.boxplot(x=df_box['funding_total_usd'], y=df_box['status'])
```

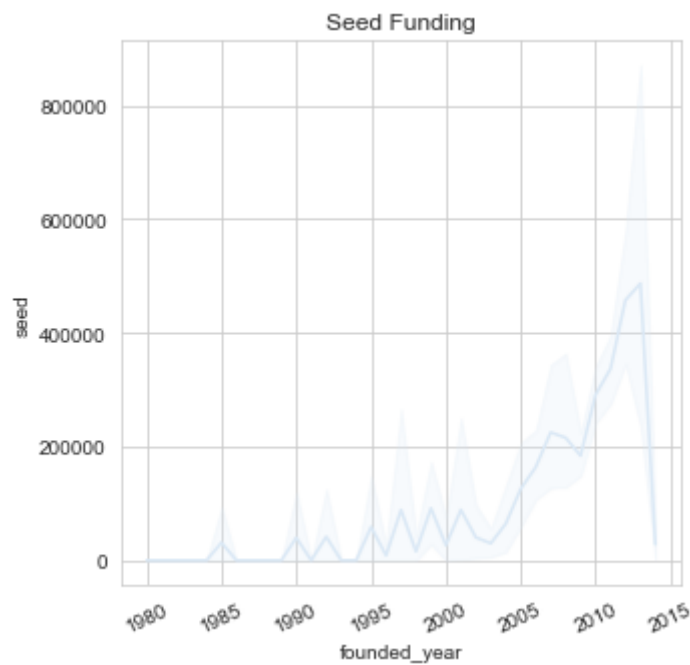

```
Out[134]: <AxesSubplot:xlabel='funding_total_usd', ylabel='status'>
```



```
In [135]: plt.figure(figsize=(5,5))
plt.title('Venture Capital Funding')
sns.lineplot(x = "founded_year", y = "venture", data = df)
plt.xticks(rotation = 25)

plt.figure(figsize=(5,5))
sns.lineplot(x = "founded_year", y = "seed", data = df)
plt.xticks(rotation = 25)
plt.title('Seed Funding')
```

```
Out[135]: Text(0.5, 1.0, 'Seed Funding')
```




```
In [136... # top 10 markets

#print(set(df['market'][df['market'].map(df['market'].value_counts()) >= 100]))
df['market'].value_counts()[1:11]
top10_mkt = list(set(df['market'][df['market'].map(df['market'].value_counts())
top10_mkt

plt.figure(figsize=(10,10))
sns.catplot(x='market',hue='status',data=df[df['market'].isin(top10_mkt)],
            kind="count",
            order = df['market'][df['market'].isin(top10_mkt)].value_counts().i
plt.xticks(rotation = 65)
```

```
Out[136]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
 [Text(0, 0, 'other'),
  Text(1, 0, 'Software'),
  Text(2, 0, 'Curated Web'),
  Text(3, 0, 'Mobile'),
  Text(4, 0, 'Enterprise Software'),
  Text(5, 0, 'Biotechnology'),
  Text(6, 0, 'Advertising'),
  Text(7, 0, 'Games'),
  Text(8, 0, 'E-Commerce'),
  Text(9, 0, 'Social Media'),
  Text(10, 0, 'Hardware + Software')])
<Figure size 720x720 with 0 Axes>
```



```
In [137]: print(df['region'].value_counts()[1:11])
top10_reg = list(set(df['region'][df['region'].map(df['region'].value_counts())
top10_reg

# fig= plt.figure(figsize=(10,5))
# fig, ax = plt.subplots(figsize=(10, 5))
# ax.set_xlabel("Region", fontsize=10)
# ax.set_ylabel("Count", fontsize=10)
# fig.suptitle("Status, Top 10 Regions", fontsize=15)

sns.catplot(x='region',hue='status',data=df[df['region'].isin(top10_reg)],
            kind="count",
            order = df['region'][df['region'].isin(top10_reg)].value_counts().i
plt.title('Status, Top 10 Regions')
plt.xticks(rotation = 65)
```

```

SF Bay Area      1250
New York City    338
Boston           260
Los Angeles      186
London           146
Seattle          143
Washington, D.C. 85
Chicago           85
Austin           84
San Diego        77


```

```
Name: region, dtype: int64
```

```

Out[137]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
 [Text(0, 0, 'other'),
  Text(1, 0, 'SF Bay Area'),
  Text(2, 0, 'New York City'),
  Text(3, 0, 'Boston'),
  Text(4, 0, 'Los Angeles'),
  Text(5, 0, 'London'),
  Text(6, 0, 'Seattle'),
  Text(7, 0, 'Washington, D.C.'),
  Text(8, 0, 'Chicago'),
  Text(9, 0, 'Austin'),
  Text(10, 0, 'San Diego')])

```




```

In [138... plt.figure(figsize=(10,5))
sns.boxplot(x=df_box['venture'][df_box['founded_year'] < 2005],
            y=df_box['status'][df_box['founded_year'] < 2005])

plt.figure(figsize=(10,5))
sns.boxplot(x=df_box['venture'][df_box['founded_year'] >= 2005],
            y=df_box['status'][df_box['founded_year'] >= 2005])

```

Out[138]: <AxesSubplot:xlabel='venture', ylabel='status'>




```
In [139... # number of funding rounds
df['funding_rounds'].value_counts()

# acquired or closed by number of funding rounds

sns.catplot(x='funding_rounds', hue='status', data=df, kind="count")
```

Out[139]: <seaborn.axisgrid.FacetGrid at 0x7fc7555bef10>





```
In [140... encoder = LabelEncoder()

df_fr = df[['funding_rounds', 'status']]
df_fr['cat'] = pd.Series(encoder.fit_transform(df['status']))

df_fr

#df_fr.groupby('status').mean()
#df_fr.groupby('funding_rounds').sum()
df_fr_group = pd.DataFrame(df_fr[['funding_rounds', 'cat']].groupby('funding_rounds').sum())

df_fr_group = df_fr_group.droplevel(axis=1, level=0).reset_index()
df_fr_group['pct_fail'] = df_fr_group['sum'] / df_fr_group['count']

display(df_fr_group)


sns.set_style("whitegrid")
sns.set_palette('Set2')
g = sns.pointplot(data=df_fr_group[:6], x="funding_rounds", y="pct_fail")
g.set(xlabel='Number of Funding Rounds',
      ylabel='Failure Rate',
      title='Percent of Failed Companies by Number of Funding Rounds')

vals = g.get_yticks()
g.set_yticklabels(['{:, .0%}'.format(x) for x in vals])

vals = g.get_xticks()
g.set_xticklabels(['{:, .0f}'.format(x) for x in vals])
```

	funding_rounds	count	sum	pct_fail
0	1.0	2793	1456	0.521303
1	2.0	1073	339	0.315937
2	3.0	552	100	0.181159
3	4.0	263	51	0.193916
4	5.0	118	20	0.169492
5	6.0	70	7	0.100000
6	7.0	32	5	0.156250
7	8.0	19	8	0.421053
8	9.0	5	0	0.000000
9	10.0	2	0	0.000000
10	11.0	2	0	0.000000
11	15.0	1	0	0.000000

```
Out[140]: [Text(0, 0, '0'),
Text(1, 0, '1'),
Text(2, 0, '2'),
Text(3, 0, '3'),
Text(4, 0, '4'),
Text(5, 0, '5')]
```



```
In [141]: df[df['funding_rounds']==8]
```

Out[141]:

	name	market	funding_total_usd	status	country_code	state_code	reg
<b>147</b>	Aggregate Knowledge	Advertising	7.043011e+11	acquired	USA	CA	SF
<b>258</b>	Appia	Advertising	4.042050e+11	acquired	USA	NC	ot
<b>505</b>	Biolex Therapeutics	Biotechnology	1.702206e+12	closed	USA	NC	ot
<b>798</b>	CipherMax	Security	1.401504e+12	closed	USA	CA	SF
<b>848</b>	Cloudant	Enterprise Software	1.082050e+11	acquired	USA	MA	Bos
<b>985</b>	Cozi Group	other	2.095057e+11	closed	USA	WA	Sea
<b>1147</b>	Dilithium Networks	Mobile	9.096070e+11	closed	USA	CA	ot
<b>1416</b>	Extreme Enterprises	Enterprise Software	1.405500e+09	closed	USA	FL	ot
<b>1523</b>	Flurry	Mobile	7.032055e+11	acquired	USA	CA	SF
<b>2295</b>	Laszlo Systems	other	3.098090e+11	acquired	USA	CA	SF
<b>2664</b>	Moblyng	Games	1.091030e+11	closed	USA	CA	SF
<b>3136</b>	PayScale	other	3.033086e+11	acquired	USA	WA	Sea
<b>3253</b>	Plextronics	Clean Technology	5.014047e+11	acquired	USA	PA	ot
<b>3927</b>	Socialtext	Enterprise Software	4.067070e+11	acquired	USA	CA	SF
<b>3941</b>	SolFocus	Clean Technology	2.101400e+12	closed	USA	CA	SF
<b>3951</b>	Solyndra	Manufacturing	1.056075e+14	closed	USA	CA	SF
<b>4602</b>	Virident Systems	other	1.402303e+12	acquired	USA	CA	SF
<b>4793</b>	Xobni	Software	4.017052e+11	acquired	USA	CA	SF
<b>4860</b>	Zappos	Curated Web	6.027050e+11	acquired	USA	NV	ot

19 rows x 33 columns

In [142...

```
df_fr.groupby('status').mean()
#df_fr.groupby('funding_rounds').sum()
```

```
Out[142]:
```

	funding_rounds	cat
status		
acquired	2.125000	0
closed	1.449648	1

```
In [143... df_time = df[['status', 'time_to_first_funding']]
df_time.groupby('status').mean()
```

```
Out[143]:
```

	time_to_first_funding
status	
acquired	3.740585
closed	1.685871

```
In [144... sns.set_style("whitegrid")
sns.set_palette('Set2')

#category_order = ["U.S.", "International"]

g = sns.catplot(x='international', hue='status', data=df, kind="count")
g.set_xticklabels(["U.S.", "International"])
g.set(xlabel='Location', ylabel='Company Count')
g.fig.suptitle("Companies by Status & Location", y=1.02);
```



```
In [145... # ## region other in train / test, need to add to df


# sns.set_style("whitegrid")
# sns.set_palette('Set2')

# #category_order = ["U.S.", "International"]
```

```
# g = sns.catplot(x='region_other',hue='status', data=df, kind="count")
# g.set_xticklabels(["Top -- Region", "Other Region"])
# g.set(xlabel='Location', ylabel='Company Count')
# g.fig.suptitle("Companies by Status & Region", y=1.02);
```


```
In [146]: # more than one funding round
df_ofr = df_box[df_box['one_funding_round']=='one']
df_mfr = df_box[df_box['one_funding_round']=='multiple']
sns.boxplot(x='venture', y='status', data=df_ofr)
```

Out[146]: <AxesSubplot:xlabel='venture', ylabel='status'>



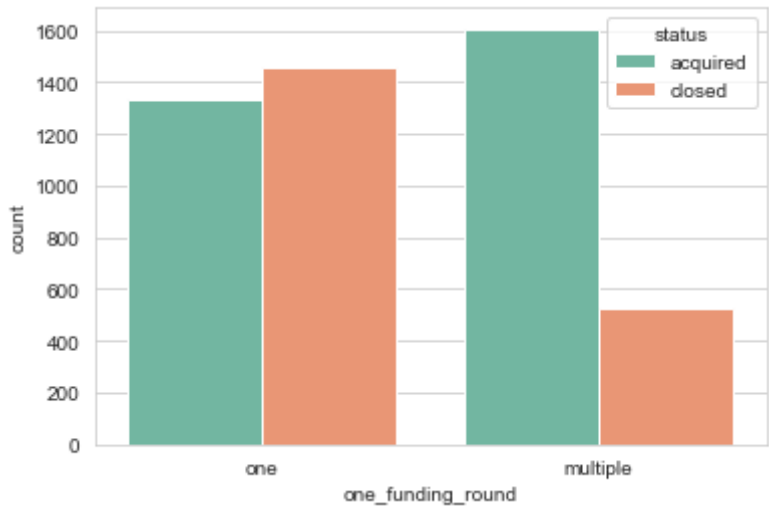
```
In [147]: sns.boxplot(x='venture', y='status', data=df_mfr)
```

Out[147]: <AxesSubplot:xlabel='venture', ylabel='status'>



```
In [148]: sns.countplot(x='one_funding_round',hue='status', data=df, orient='v')
```

Out[148]: <AxesSubplot:xlabel='one\_funding\_round', ylabel='count'>



```
In [149... ### need something with funding_total_usd, time to first funding
```